Structuring the empirical chapters

A skeleton for the data-collection / results and discussion / evaluation sections, with the figure and number for each claim already attached. Figure numbers refer to `docs/report_figures/`.

The two chapters can be combined — the structure below works either as two chapters or as one chapter with results and interpretation interleaved per axis. **If you combine them, do it per axis** (tolerance: result then interpretation; cost: result then interpretation), not by putting all results first and all discussion second. Mixed-order combining is where these chapters usually become hard to follow.

Chapter A — Experimental method and data collection

The job of this chapter is to make the results *believable before they are read*. Everything here is about controls.

A.1 What is held constant, and why it has to be

The central methodological claim: comparing erasure codes is meaningless unless storage overhead is fixed, because any code tolerates more erasure if allowed more space. Every scheme runs at **$n \approx 8192$, $k = 4096$, ~2× overhead** (2D-RS at 2.06×, within 3%).

State plainly that an earlier iteration had 2D-RS at 4× overhead, and that comparing a 4×-overhead code against 2×-overhead codes at equal erasure rate flattered it. Describing a correction you made is a strength, not an admission — it shows the control was deliberate.

A.2 The fairness controls

Four, in increasing order of strictness:

1. **One field.** Real BLS12-381 scalar field, not a toy field.
2. **One blob.** Real application bytes packed once via EIP-4844 rules; every scheme encodes that same blob, not seven random ones.
3. **Identical erasure positions.** Where codeword lengths match — 1D-RS and all five BC rows are $n = 8192$ — trials corrupt the *same positions*, not merely the same rate. 2D-RS has $n = 8448$ and so draws its own masks at the same rate. Say this explicitly; it is the strictest control available and a reader will wonder.
4. **Success is measured in bytes.** A trial counts only if the recovery round-trips to the exact original application bytes.

A.3 Why byte-level verification, not the decoder's own flag

This deserves its own subsection, and **figure 13**. An earlier decoder returned zeros for values it had not solved and reported success: across 40 trials it claimed success every time and was wrong on 31. Had the harness trusted the `remaining` flag, it would have reported 100% success at every erasure rate.

This is the strongest methodological point in the project. It shows the validation caught a fault that would have invalidated every number, and it justifies the expense of byte-level checking.

A.4 Trial allocation and the threshold definition

- Trial counts vary by scheme and rate — more near a transition where outcomes genuinely vary, fewer where they do not. Give the rule, and note that 95% Wilson intervals are in the JSON.
- **Define the threshold explicitly**: the erasure rate at which the success rate crosses 50%, linearly interpolated between the two swept rates that bracket it. Every tolerance number in the report depends on this definition, so it must be stated once, early, and not changed.

A.5 Reproducibility

Measurement writes JSON; every figure and table is computed from that JSON rather than transcribed. Figure generation is deterministic. State the pipeline — `benchmark_1D_2D_BC.py` → `benchmark_results.json` → `make_figures.py` — and that 520 trials underlie the sweep.

Chapter B — Results

One section per axis from the research question. Report numbers; save interpretation for Chapter C (or the paired subsection, if combining).

B.1 Fault tolerance — figures 1, 2

scheme	measured threshold
1D-RS	0.505
2D-RS	0.397
BC ω512 λ2	0.350
BC ω1024 λ2	0.341
BC ω2048 λ2	0.340
BC ω1024 λ3	0.275
BC ω1024 λ4	0.210

Two findings:

- **1D-RS reaches 0.505 against a theoretical maximum of 0.50** — what being MDS means, and a sanity check that the harness is sound.
- **The block-circulant threshold is set by λ alone.** The three λ=2 rows land at 0.340, 0.341 and 0.350 while μ ranges over 2, 4, 8 and ω over 512 to 2048. Against the $1/(1+\lambda)$ prediction of 0.333, 0.250, 0.200, the measurements are 0.34, 0.28, 0.21. Figure 2 is theory and experiment agreeing — lead with it.

B.2 Computational cost — figures 4, 6

Report encode and decode separately; they are not interchangeable. Decode at 20% erasure: 1D-RS 20.4 s, 2D-RS 2.7 s, BC rows 1.5–5.5 s. Encode runs the other way: 1D-RS 0.23 s, BC ω2048 15.7 s.

Note that all timings are single-threaded Python over `galois`, so they are **relative** measurements, not production figures. Say so once here rather than defending it later.

B.3 Bandwidth — **figure 8**

Surviving data needed to rebuild one blob, as a multiple of blob size: 1D-RS 0.99× (the floor), everything else 1.24–1.58×.

B.4 Network behaviour — **figures 9, 10**

The simulation's distinguishing feature: in the benchmark the erasure mask is an *input you invent*; here it is an *output* — whatever failed to arrive before the deadline. Make that the opening sentence of the section.

- Realised erasure ≈ p^4 under custody replication 4 (figure 9), identical across all six schemes.
- Schemes fail in the same order as the coding benchmark predicted (figure 10), via an entirely different mechanism. **This is the most valuable result in the project** — two independent methods agreeing is far stronger evidence than either alone. Do not bury it.

B.5 Scalability — **figure 12**

Realised erasure against peer count, against the exact hypergeometric probability computed from the model rather than fitted. Measurement lies on prediction. Conclusion: what protects a cell is the *fraction* of the network custodying it, so peer count cannot be varied without specifying replication.

Chapter C — Discussion and evaluation

C.1 Guarantees versus behaviour — **figure 3**

2D-RS tolerates 0.397 against a worst-case guarantee of 0.098 — four times its own bound. Not a contradiction: a minimum-distance guarantee holds against *every* pattern including the adversarial worst case, while these trials draw patterns at random.

This distinction is worth a paragraph of its own, because it is the most likely place for a reader to think they have found an error.

C.2 Why the costs differ — **figure 5**

Raw decode time rises with erasure rate for everything, which mostly says there is more to repair. Divide it out and the structure appears: 1D-RS flat at 10–13 ms per erasure (cost linear in erasures, one long division's worth each); BC rows at 1–3.5 ms; 2D-RS falling from 8 ms to 1 ms, revealing a large fixed cost that amortises — which is why it looks expensive at 2% erasure and cheap at 30%.

C.3 The trade-off — **figure 7**

The synthesis. Only three of seven schemes are on the cost-versus-tolerance frontier. BC ω512 λ2 is the cheapest scheme still surviving a third of the codeword; 2D-RS buys ∼5 points more tolerance for roughly double the decode; 1D-RS buys the maximum and pays an order of magnitude.

Call out explicitly that **BC ω1024 λ4 is dominated on both axes** by BC ω512 λ2 — a parameter choice with nothing to recommend it. Negative results about your own parameter space are credible and worth stating.

C.4 The asymmetry that changes the recommendation — figure 6

Encode is paid once by the proposer; decode is paid by every reconstructing node. 1D-RS encodes in 0.23 s and decodes in 20 s; BC ω2048 spends 15.7 s encoding to decode in 5.5 s. For data-availability sampling the second shape is the more useful one — an argument for block-circulant the tolerance numbers alone do not make.

C.5 The result that reframes the rest — figure 11

Every scheme gathers in about a second; decode spans three orders of magnitude above that and misses the 4-second sampling deadline. **Bandwidth was the expected constraint and is not close.** Decode arithmetic is the binding one.

This has a consequence worth stating: it reorders the research question's own priorities. A scheme's erasure tolerance matters less than expected, because custody replication absorbs most peer failure (§B.4) and decode speed dominates the time budget.

C.6 Threats to validity

Be first to raise these. A reader who finds them unprompted discounts everything else.

- **Random, not adversarial, erasures.** Patterns are drawn i.i.d.; a real adversary chooses the worst pattern. Measured tolerances are therefore upper bounds on adversarial performance — directly relevant to 2D-RS and the λ=2 rows, which exceed their worst-case guarantees.
- **Single-threaded Python.** Relative comparisons hold; absolute seconds would not survive an optimised implementation.
- **The block-circulant rows carry a ~2.3× implementation penalty** that the Reed-Solomon rows do not. Defining the BC local code over roots of unity (to make it FFT-capable) means solving over 255-bit field elements instead of small integers, and that cost is paid even with the transform switched off. Measured in `decoder_version_results.json`. This narrows the apparent 1D-RS gap and moves the 2D-RS crossover; it does not touch any tolerance result. Raise it yourself in §B.2 — it is the kind of thing a reader finds and then distrusts everything else.
- **One blob of application data.** Erasure patterns are reseeded per trial, but the payload is fixed.
- **Block-circulant is measured single-phase.** Phase 2 does not implement the paper's auxiliary construction, so BC thresholds are what phase 1 alone achieves — a limit of this implementation, not of the code. Say which.
- **One network topology**, where the brief asks for several.
- **Simulated network.** Latency and bandwidth are modelled, not measured on real hardware.

C.7 Limitations and further work

Three concrete items, each with a named fix:

1. Implement Algorithm 1's auxiliary GRS construction for phase 2 — would raise the correction radius from ρ toward 2ρ and drop the cost from cubic to quadratic.
2. Replace the NTT path's polynomial division with Newton iteration — the transform is only 12% of that decode; ~80% is a schoolbook division.
3. Add further network topologies.

Two things to settle before writing

The NTT decoder's status — settled. The roots-of-unity version stays; nothing is being reverted and no re-run is needed. The decode numbers in the report are the ones to write about. The ~2.3× penalty this carries for the block-circulant rows is measured, stated in the report, and listed under threats to validity above — handle it there rather than by adjusting numbers.

Where the NTT investigation lives. It is a genuine negative result — implemented, measured, 7–23× slower, with the cause identified — but it is about *implementation*, not about the coding schemes. It belongs in an implementation appendix or a short subsection of C.7, not in the empirical results. Its four figures are in `docs/figures/` and deliberately left out of `docs/report_figures/`.

Figure budget

Thirteen figures is a lot for one chapter. If you need to cut, in order of what goes first:

keep	if space
07 (the trade-off — if you keep one, keep this)	03
01, 02 (tolerance, and theory agreeing)	05
11 (the bottleneck)	08
10 (the two methods agreeing)	12
04 (decode cost)	13 (move to methodology)
09 (p^4)	06